

Shared memory computing in R

Course: Parallel computing in R

Overview

1. Execution types in R
2. Using optimized libraries
3. The parallel package (shared)
4. The foreach package (shared, distributed)

Execution types in R

- Default R only uses serial computing
- Shared memory parallel computing is using...
 - Thread, Pthread (posix thread)
 - Sockets
 - Rarely OpenMP (Open Multi-Processing, multithreading)
- Distributed memory computing
 - Poor scaling of R above one node

Using optimized libraries

Faster linear algebra functions in R

Basic Linear Algebra Subroutines (BLAS)

Provide standard building blocks for performing basic vector and matrix operations

BLAS

- Usually installed by default
- Default **BLAS** is not optimized

Installation

1. In *UBUNTU* install **OpenBLAS** via

```
$ sudo apt-get install libopenblas-base
```

2. in *OSX* install **OpenBLAS** via

```
$ brew install openblas
```

3. In windows, please check out *Microsoft R* <https://mran.microsoft.com>

Installation of OpenBLAS

1. Before installation

```
> La_library()  
[1] "/usr/lib/x86_64-linux-gnu/lapack/liblapack.so.3.9.0"  
> extSoftVersion()["BLAS"]  
  
BLAS  
"/usr/lib/x86_64-linux-gnu/blas/libblas.so.3.9.0"
```

2. After installation (**POSIX threads available**)

```
> La_library()  
[1] "/usr/lib/x86_64-linux-gnu/openblas-pthread/liblapack.so.3"  
> extSoftVersion()["BLAS"]  
  
BLAS  
"/usr/lib/x86_64-linux-gnu/openblas-pthread/libblas.so.3"
```

Linear algebra computing comparison

- simulate a matrix of 1 million observations by n predictors and generate an outcome y .
- Compute the least squares estimates of the linear regression coefficients when regressing the response y on the predictor matrix X

Function

```
simLR <- function(n) {  
  x <- matrix(rnorm(1e6 * n), 1e6, n)  
  b <- rnorm(n)  
  y <- drop(x %*% b) + rnorm(n)  
  b <- solve(crossprod(x), crossprod(x, y))  
}
```


Testing with different BLAS

```
start_time <- Sys.time()  
simLR(100)  
print(Sys.time() - start_time)
```

Without optimized BLAS

Execution time of 16.81325 secs

With optimized BLAS

Execution time of 7.17849 secs

The parallel package

Parallel package

1. Installed by default

```
library(parallel)
```

2. Contains many functions
3. Can detect how many cores you have

```
detectCores(logical = TRUE/FALSE)
```

Hyperthreading

Process where a **CPU** splits each of its physical cores into virtual cores, which are known as threads

LOGICAL = TRUE detects the hyperthreaded cores

Types of parallel calculations

Forking

- Faster than sockets.
- Because it copies the existing version of R, your entire workspace exists in each process.
- Trivially easy to implement.
- Only works on POSIX systems (Mac, Linux, Unix, BSD) and not Windows.
- Because processes are duplicates, it can cause issues specifically with random number generation

Types of parallel calculations

Socket

- Works on any system (including Windows).
- Each process on each node is unique so it can't cross-contaminate
- Each process is unique so it will be slower
- Things such as package loading need to be done in each process separately. Variables defined on your main version of R don't exist on each core unless explicitly placed there.
- More complicated to implement.

Forking parallel example

List of functions

```
mcLapply(X, FUN, ...)  
mcmapply(X, FUN, ...)
```

Example

```
library(parallel)  
no_cores <- detectCores() - 1  
mcLapply(1:100, FUN,  
         mc.cores = no_cores)
```

Decrease *no_cores* by one for not interfering with master thread

With small tasks, the overhead of scheduling the task and returning the result can be greater than the time to execute the task itself

Socket parallel example

List of functions

```
parLapply(cluster, X, FUN)  
parSapply(cluster, X, FUN)
```

Decrease *no_cores* by one for not interfering with master thread

Example

```
library(parallel)  
no_cores <- detectCores() - 1  
cl <- makeCluster(no_cores)  
parLapply(cl, 1:100, FUN)  
stopCluster(cl)
```

Sockets: Passing objects

- If libraries are used you need to load the package within the process

```
clusterEvalQ(cl, library([LIBRARY NAME]))
```

- Variables are not present within the parallelisation and must be passed.

```
# Move variable x to the cluster  
clusterExport(cl, "x")  
# Register the variable x on the cluster  
clusterEvalQ(cl, x)
```


The foreach package

Works for serial, shared and distributed memory computing

Serial Foreach example

- We have a list of vectors (Default dataset of co2 concentrations in ppm 1959 to 1997)

```
# Create a list with values from each year
x <- split(co2, ceiling(seq_along(co2)/12))
foreach(i = x) %do%
  mean(i)
```

Using Foreach package in parallel

1. Must be installed

```
library(parallel)  
library(foreach)  
library(doParallel)
```

2. Similar use as **OpenMP** <https://en.wikipedia.org/wiki/OpenMP>

3. More information at <https://cran.r-project.org/web/packages/foreach/foreach.pdf>

How about parallel calculations: doParallel package

1. Packages needed

```
library(parallel)  
library(foreach)  
library(doParallel)
```

2. Shows how many parallel processes are available

```
getDoParWorkers()
```

3. Register the number of processes

```
registerDoParallel()
```

4. Substitute %do% with %dopar%

Parallel shared memory example

```
library(parallel)
library(foreach)
library(doParallel)
x <- split(co2, ceiling(seq_along(co2)/12))
registerDoParallel()
foreach(i = x) %do%
  mean(i)
```

1. Adding extra libraries
2. Register the number of processes (Default: all)
3. Switch *do* with *dopar*

Defining the number of parallel processes

1. Set the number of processes

```
nproc <- makeCluster(<number of processes>)
```

2. Register the number of processes

```
registerDoParallel(nproc)
```

3. End all processes

```
stopCluster(nproc)
```

Useful parameters

Process the tasks results as they are generated

```
foreach(i = 1:100,  
  .combine='[type]') %do%  
  FUN(i)
```

Type	Process
c	returns an array
cbind	returns a column matrix
rbind	returns a row matrix
+	summarize the results
*	Multiply the results

Working examples

1. Equivalent to **lapply(x, mean)**

```
x <- split(co2, ceiling(seq_along(co2)/12))  
foreach(i = x) %do%  
  mean(i)
```

2. Return array of means

```
x <- split(co2, ceiling(seq_along(co2)/12))  
foreach(i = x, .combine='c') %do%  
  mean(i)
```

3. replace *do* with *dopar* for parallelisation

Distributed computing

How to use foreach for distributed computing

- Can be used for distributed computing using *doMPI* library
- R seldomly scales above one node on a supercomputer

```
library(doMPI)
cl <- startMPIcluster()
# You define N processes when allocating your job
registerDoMPI(cl)
x <- split(co2, ceiling(seq_along(co2)/12))
foreach(i = x) %dopar%
  mean(i)
closeCluster(cl)
mpi.quit()
```

